

ML-Lab1

November 9, 2021

1 Lab 1: Decision Trees

Parand Mohri

```
[149]: import pandas as pd
from sklearn import tree
from sklearn.metrics import accuracy_score
from sklearn.model_selection import train_test_split
import matplotlib.pyplot as plt
import numpy as np
```

1) Diabetes Data

```
[150]: data = pd.read_csv('/Users/macbook/Downloads/diabetes.csv')
Y = data['class']
X = data.drop(['class'],axis=1)
```

Deviding the data into test and train and training one-level decision tree:

```
[151]: clf = tree.DecisionTreeClassifier(criterion = 'entropy',
max_depth = 1)
X_train, X_test, Y_train, Y_test = train_test_split(X, Y,
test_size=0.34, random_state=10)
clf = clf.fit(X_train, Y_train)
```

Calculating the acuuracy of the Train data:

```
[152]: Yp = clf.predict(X_train)
acc = accuracy_score(Y_train, Yp)
print(acc)
```

0.7648221343873518

Calculating the acuuracy of the Test data:

```
[153]: Yp = clf.predict(X_test)
acc = accuracy_score(Y_test, Yp)
print(acc)
```

0.7213740458015268

With one-level decision tree we got the above accuracies for train and test data. As you can see train data has better accuracy than test data and that is not a surprise because the tree was made by train data so it's more accurate for that data.

Dividing the data into test and train and training multi-level decision tree:

```
[154]: clf = tree.DecisionTreeClassifier(criterion = 'entropy',
        max_depth = None)
X_train, X_test, Y_train, Y_test = train_test_split(X, Y,
        test_size=0.34, random_state=10)
clf = clf.fit(X_train, Y_train)
```

Calculating the accuracy of the Train data:

```
[155]: Yp = clf.predict(X_train)
acc = accuracy_score(Y_train, Yp)
print(acc)
```

1.0

Calculating the accuracy of the Test data:

```
[156]: Yp = clf.predict(X_test)
acc = accuracy_score(Y_test, Yp)
print(acc)
```

0.7251908396946565

With multi-level decision tree we got the above accuracies for train and test data. The accuracy for training data is 1 because the tree doesn't have any limits for max depth so it makes new branches for every data if needed so it makes the right decision for all the data in train data, this is not true for test data because it can be some cases in the test data that the decision tree doesn't know about.

Putting min_samples_leaf to the size of the datasets and dividing the data into test and train and training the decision tree:

```
[157]: clf = tree.DecisionTreeClassifier(criterion = 'entropy',
        max_depth = None, min_samples_leaf = 768)
X_train, X_test, Y_train, Y_test = train_test_split(X, Y,
        test_size=0.34, random_state=10)
clf = clf.fit(X_train, Y_train)
```

Calculating the accuracy of the Train data:

```
[158]: Yp = clf.predict(X_train)
acc = accuracy_score(Y_train, Yp)
print(acc)
```

0.6561264822134387

Calculating the accuracy of the Test data:

```
[159]: Yp = clf.predict(X_test)
acc = accuracy_score(Y_test, Yp)
print(acc)
```

0.6412213740458015

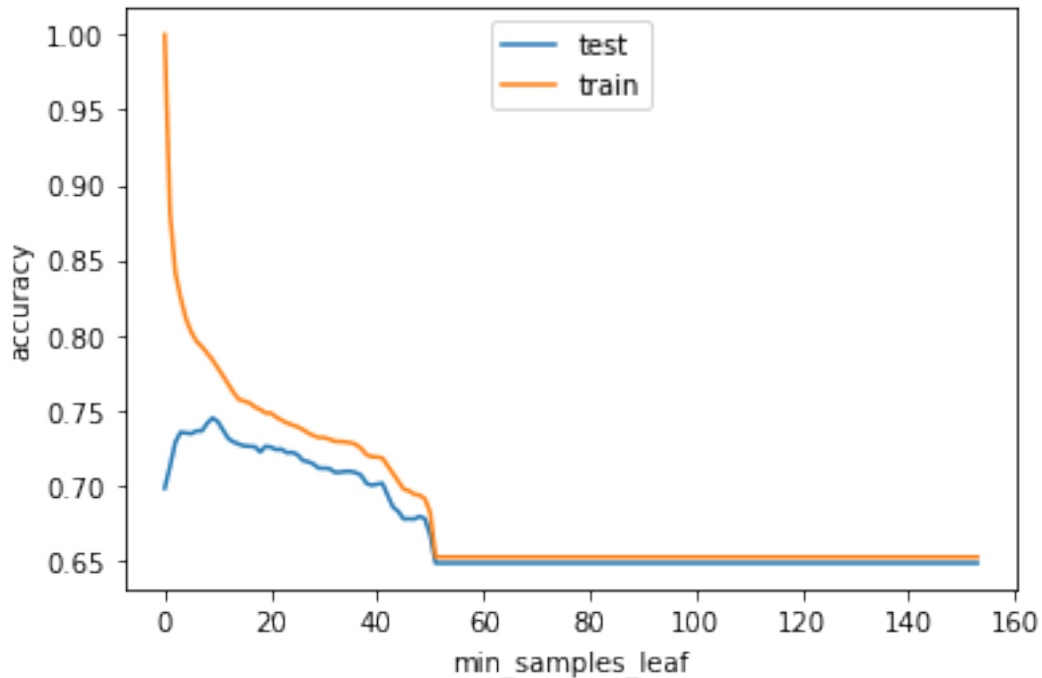
With multi-level decision tree with `min_samples_leaf` equal to the size of the datasets we got the above accuracies for train and test data. The accuracies here are less than what we have for multi-level and one-level decision trees, the reason is that we are stopping when the leaf nodes get less than 768 training instances and that gonna leave us with only the root node because that is the size of our dataset. so here our knowledge about test and training data is not that different and we can even have higher accuracy for our test data than train data as you can see it happened here as well.

Plot the accuracy rates based on the training set and hold-out validation for `min_samples_leaf` from 1 to the size of the datasets with step of 5: Here im generating 50 trees with different test and train datas and making the graph with average data

```
[195]: def accuracy(X_train, X_test, Y_train, Y_test):
    test = []
    train = []
    for i in range(1,768,5):
        clf = tree.DecisionTreeClassifier(criterion = 'entropy',
                                          max_depth = None, min_samples_leaf = i)
        clf = clf.fit(X_train, Y_train)
        Yp_test = clf.predict(X_test)
        Yp_train = clf.predict(X_train)
        test.append(accuracy_score(Y_test, Yp_test))
        train.append(accuracy_score(Y_train, Yp_train))
    return test , train

acTest = [0] * round(768/5)
acTrain = [0] * round(768/5)
for i in range(50):
    X_train, X_test, Y_train, Y_test = train_test_split(X, Y,
                                                         test_size=0.34)
    acTest = np.add(acTest, accuracy(X_train, X_test, Y_train, Y_test)[0])
    acTrain = np.add(acTrain, accuracy(X_train, X_test, Y_train, Y_test)[1])
acTest = [x / 50 for x in acTest]
acTrain = [x / 50 for x in acTrain]
```

```
[197]: plt.plot(acTest,label = 'test')
plt.plot(acTrain,label = 'train')
plt.ylabel('accuracy')
plt.xlabel('min_samples_leaf')
plt.legend(loc ='upper center')
plt.show()
```



this is the resulting graph. It is really make sense because as we start our min-sample-leaf is low and we have a multi-level tree so its going to predict the train data much better than test data but as min-sample-leaf grow they get closer and closer together and after some time we see that the accuracy doesnt change for neither of them. and we can see the overfitting of training data.

2) Glass Data

```
[14]: data = pd.read_csv('/Users/macbook/Downloads/glass.csv')
      Y = data['class']
      X = data.drop(['class'],axis=1)
```

Deviding the data into test and train and training one-level decision tree:

```
[15]: clf = tree.DecisionTreeClassifier(criterion = 'entropy',
      max_depth = 1)
      X_train, X_test, Y_train, Y_test = train_test_split(X, Y,
      test_size=0.34, random_state=10)
      clf = clf.fit(X_train, Y_train)
```

Calculating the acuuracy of the Train data:

```
[16]: Yp = clf.predict(X_train)
      acc = accuracy_score(Y_train, Yp)
      print(acc)
```

0.46099290780141844

Calculating the accuracy of the Test data:

```
[17]: Yp = clf.predict(X_test)
      acc = accuracy_score(Y_test, Yp)
      print(acc)
```

0.4246575342465753

With one-level decision tree we got the above accuracies for train and test data. As you can see train data has better accuracy than test data and that is not a surprise because the tree was made by train data so it's more accurate for that data.

Dividing the data into test and train and training multi-level decision tree:

```
[18]: clf = tree.DecisionTreeClassifier(criterion = 'entropy',
      max_depth = None)
      X_train, X_test, Y_train, Y_test = train_test_split(X, Y,
      test_size=0.70, random_state=10)
      clf = clf.fit(X_train, Y_train)
```

Calculating the accuracy of the Train data:

```
[19]: Yp = clf.predict(X_train)
      acc = accuracy_score(Y_train, Yp)
      print(acc)
```

1.0

Calculating the accuracy of the Test data:

```
[20]: Yp = clf.predict(X_test)
      acc = accuracy_score(Y_test, Yp)
      print(acc)
```

0.54

With multi-level decision tree we got the above accuracies for train and test data. The accuracy for training data is 1 because the tree doesn't have any limits for max depth so it makes new branches for every data if needed so it makes the right decision for all the data in train data, this is not true for test data because it can be some cases in the test data that the decision tree doesn't know about.

Putting `min_samples_leaf` to the size of the datasets and dividing the data into test and train and training the decision tree:

```
[21]: clf = tree.DecisionTreeClassifier(criterion = 'entropy',
      max_depth = None, min_samples_leaf = 214)
      X_train, X_test, Y_train, Y_test = train_test_split(X, Y,
      test_size=0.34, random_state=10)
      clf = clf.fit(X_train, Y_train)
```

Calculating the accuracy of the Train data:

```
[22]: Yp = clf.predict(X_train)
      acc = accuracy_score(Y_train, Yp)
```

```
print(acc)
```

0.3546099290780142

Calculating the accuracy of the Test data:

```
[23]: Yp = clf.predict(X_test)
      acc = accuracy_score(Y_test, Yp)
      print(acc)
```

0.3561643835616438

With multi-level decision tree with `min_samples_leaf` equal to the size of the datasets we got the above accuracies for train and test data. The accuracies here are less than what we have for multi-level and one-level decision trees, the reason is that we are stopping when the leaf nodes get less than 214 training instances and that gonna leave us with only the root node because that is the size of our dataset. So here our knowledge about test and training data is not that different and we can even have higher accuracy for our test data than train data as you can see it happen here as well.

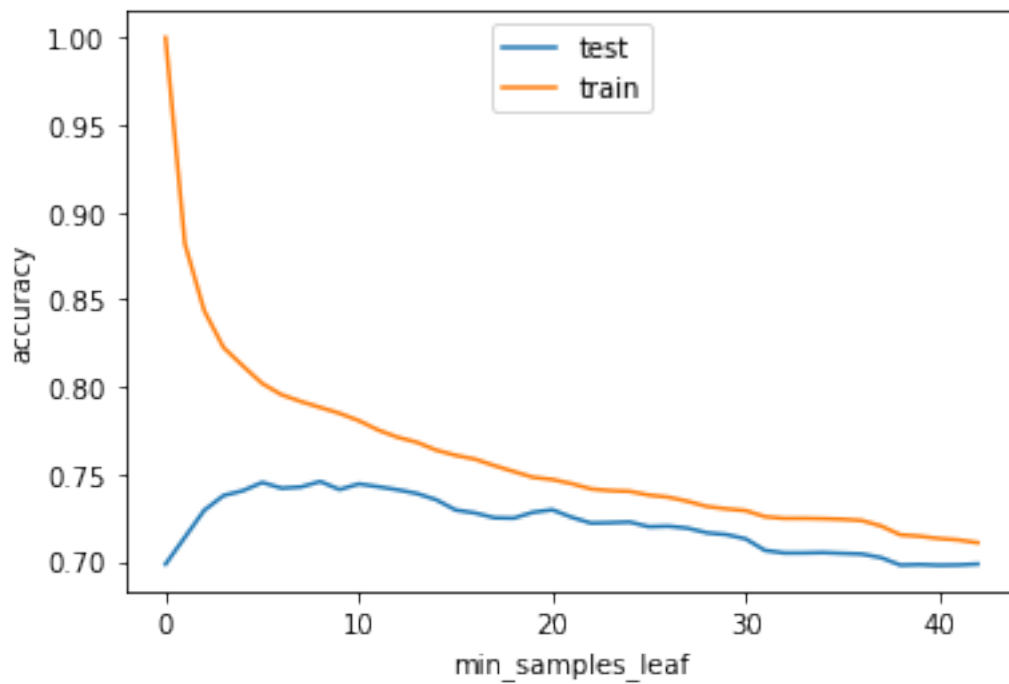
Plot the accuracy rates based on the training set and hold-out validation for `min_samples_leaf` from 1 to the size of the datasets with step of 5: Here I'm generating 50 trees with different test and train data and making the graph with average data

```
[202]: def accuracy(X_train, X_test, Y_train, Y_test):
        test = []
        train = []
        for i in range(1,214,5):
            clf = tree.DecisionTreeClassifier(criterion = 'entropy',
                                              max_depth = None, min_samples_leaf = i)
            clf = clf.fit(X_train, Y_train)
            Yp_test = clf.predict(X_test)
            Yp_train = clf.predict(X_train)
            test.append(accuracy_score(Y_test, Yp_test))
            train.append(accuracy_score(Y_train, Yp_train))
        return test , train

acTest = [0] * round(214/5)
acTrain = [0] * round(214/5)
for i in range(50):
    X_train, X_test, Y_train, Y_test = train_test_split(X, Y,
                                                         test_size=0.34)
    acTest = np.add(acTest, accuracy(X_train, X_test, Y_train, Y_test)[0])
    acTrain = np.add(acTrain, accuracy(X_train, X_test, Y_train, Y_test)[1])
acTest = [x / 50 for x in acTest]
acTrain = [x / 50 for x in acTrain]
```

```
[203]: plt.plot(acTest,label = 'test')
      plt.plot(acTrain,label = 'train')
      plt.ylabel('accuracy')
      plt.xlabel('min_samples_leaf')
```

```
plt.legend(loc='upper center')  
plt.show()
```



this is the resulting graph. It is really make sence because as we start our min-sample-leaf is low and we have a multi-level tree so its going to predict the train data much better than test data but as min-sample-leaf grow they get closer and closer together. and we can see the overfitting of training data.